



TEXAS A&M UNIVERSITY

Engineering

An introduction on building deep models

CSCE-638 Natural Language Processing

Yaochen Xie

- Goal: implement the following methods to classify the sentiment of movie review texts
 - Naive Bayes
 - Perceptron
 - Hierarchical Attention Network (optional)

- Pipeline of implementation
 - Data (pre-)processing: write functions to convert documents into features and labels (+/-). Features for Naive Bayes and Perceptron are based on word frequency/presence.
 - Learning: estimate or optimize parameters given the processed data, in `model.train()`.
 - Inference/prediction: compute predicted class given a new data sample, in `model.classify()`.

- Hierarchical Attention Network (optional)
 - Each document is initially represented as a tensor/array of shape `[1, num_sentence (20), num_words (50)]`. Inputs:
 - i. Word tokens: each word is represented by an integer, i.e., the index of the word in the vocabulary.
 - ii. GloVe embedding: each word is represented by a vector.
 - iii. BERT/ELMO: each word is represented by a vector from pre-trained deep language models.

- Word token: given a vocabulary (dictionary) of size K , each word is an integer from $[0, K-1]$.
- To input to the GRU/attention operator, we need all words to be represented by vectors \rightarrow the embedding layer:
 - Maps word tokens into word embeddings. E.g.,
`torch.nn.Embedding(...)`, `tf.keras.layers.Embedding(...)`,
`tf.nn.embedding_lookup(...)`.

- BERT and ELMO implementation:
 - Hugging-face (<https://github.com/huggingface/transformers>)
 - Works with both PyTorch and TensorFlow

```
>>> from transformers import AutoTokenizer, AutoModel

>>> tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
>>> model = AutoModel.from_pretrained("bert-base-uncased")

>>> inputs = tokenizer("Hello world!", return_tensors="pt")
>>> outputs = model(**inputs)
```

- Recap on the attention operation:
 - Intuition: given a query, aggregate information among candidates by **weighted sum** based on similarity/relevance between each candidate and the query.
 - Query: q_1, \dots, q_n ; Key: k_1, \dots, k_m ; Value: v_1, \dots, v_m

$$o_j = \frac{1}{C} \sum_{i=1}^m f(q_j, k_i) g(v_i)$$

- HAN: **query** - a learned vector, **value and key** - from word embedding.

- Attention operations:
 - Self-attention v.s. attention with learned query
 - Additional resources:
 - <http://people.tamu.edu/~sji/classes/attn-slides.pdf>
 - <http://people.tamu.edu/~sji/classes/attn.pdf>

- An example of the attention implementation (pseudo-code):

```
x = concat([GRU_left(x), GRU_right(x)], dim=-1) # [B, S, W, D]
query = Parameter(shape=[1, D], trainable=True) # model parameters
key = x.view(B*S, W, D); value = Linear(x.view(B*S, W, D))
att_score = batch_matrix_mul(query, key.transpose([0,2,1])) #[B*S, 1, W]
att_score = Softmax(att_score, dim=1) # or any other normalization
out = batch_matrix_mul(att_score, value) #[B*S, 1, D]
out = out.view(B, S, D) # summary vector of a document
```

- Check details in the paper!



TEXAS A&M UNIVERSITY
Engineering

Building Deep Learning Environments

Spinning a GPU for you Deep Models



- Linux GPU server if applicable
- Google Colab: <https://colab.research.google.com/>
 - *Jupyter-like platform*
 - *Pros: Free GPUs, easy setup ; Con: limited use time*
- AWS/Azure DL instances
 - *Similar to GPU server, but not free after trial*
- Your Windows laptop or desktop equipped with Nvidia GPU

- Virtual environment for specific projects
 - Anaconda (recommended, <https://www.anaconda.com/>)
 - venv (<https://docs.python.org/3/library/venv.html>)
- Detached training: tmux, screen, etc.
- Installing DL frameworks
 - Tensorflow: <https://www.tensorflow.org/install>
 - PyTorch: <https://pytorch.org/get-started/locally/>



TEXAS A&M UNIVERSITY

Engineering

Building Your Neural Network: An example with PyTorch

Pipeline of Building a Deep Model



- Things you need in your code
(https://pytorch.org/tutorials/beginner/basics/quickstart_tutorial.html)
 - Dataset and dataloader
 - Your neural network model (class)
 - Training objective (loss function)
 - Optimizer, optimizing strategies



Dataset and Dataloader

- Dataset: a list of prepared data samples, or a indexable instance
 - E.g., `[(features_1, label_1), ..., (features_N, label_N)]`
 - `“class CustomDataset(Dataset)”`
- DataLoader: generates a batch of data at each iteration
 - `train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)`
 - `for batch in train_dataloader:`
`(perform training with batch)`

https://pytorch.org/tutorials/beginner/basics/data_tutorial.html



Your neural network model

- Write classes based on parent class “`torch.nn.Module`” and init parent class at the very beginning,
- Define operators and parameters in “`__init__`”, use “`torch.nn.ModuleList`” instead regular list,
- Write operations with defined operators in “`forward`”
- https://pytorch.org/tutorials/beginner/basics/buildmodel_tutorial.html

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)
```


Training

- Select hyper-parameters: learning rate, batch size, epochs, etc.
- Define dataloader, model, loss function, and optimizer
- Set model to training mode and move to GPU
- https://pytorch.org/tutorials/beginner/basics/optimization_tutorial.html

```
model.train(); model.to(device); X.to(device); y.to(device)
```

```
def train_loop(dataloader, model, loss_fn, optimizer):  
    size = len(dataloader.dataset)  
    for batch, (X, y) in enumerate(dataloader):  
        # Compute prediction and loss  
        pred = model(X)  
        loss = loss_fn(pred, y)  
  
        # Backpropagation  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()  
  
    if batch % 100 == 0:  
        loss, current = loss.item(), batch * len(X)  
        print(f"loss: {loss:>7f}  [{current:>5d}/{size:>5d}]")
```



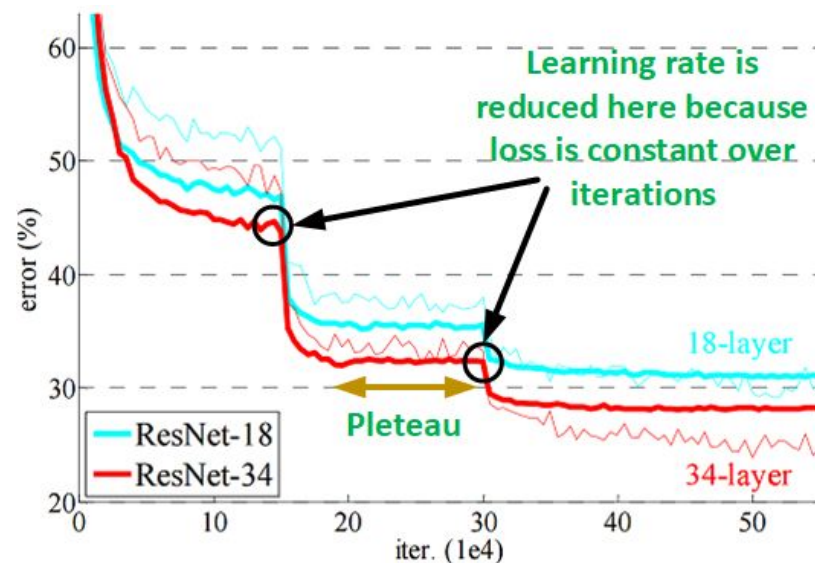
TEXAS A&M UNIVERSITY
Engineering

Some Training Tricks

- Two most common optimizer: SGD v.s. Adam
 - Stochastic gradient descent (SGD):
 - Simple yet effective, the most stable optimizer
 - Works in most case
 - Adam:
 - Use adaptive learning rate for different parameters
 - Converge faster but is unstable sometime
- <https://ruder.io/optimizing-gradient-descent/>

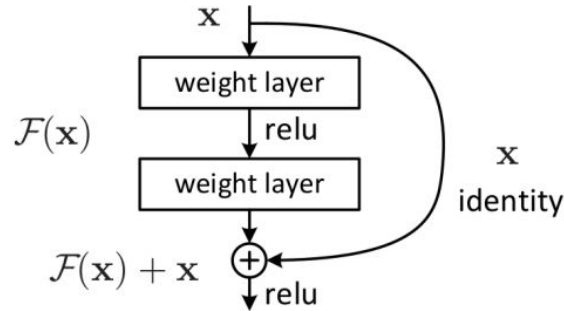
Learning Rate Scheduler

- Reduce learning rate after a certain number of epochs, or change learning rate over epochs.
- Commonly used scheduler:
 - ReduceOnPlateau, Step, Exponential, Cosine
- <https://towardsdatascience.com/learning-rate-schedules-and-adaptive-learning-rate-methods-for-deep-learning-2c8f433990d1>
- Also search for **learning rate warmup**.

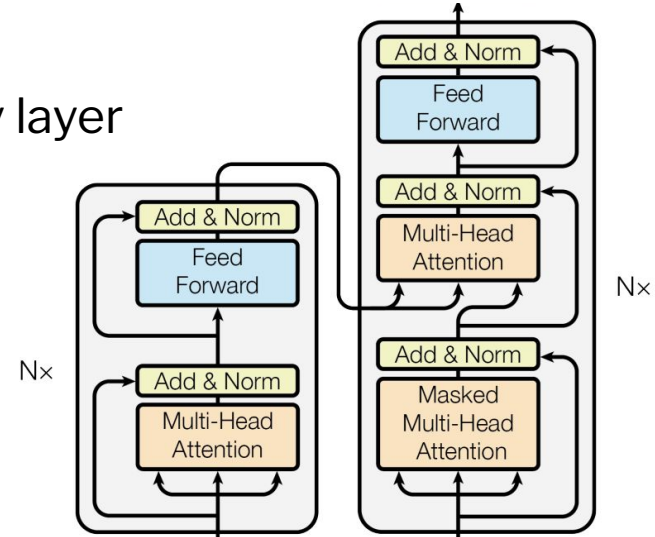


Tricks for Very Deep Models

- Skip/residual connection
 - Combines the input and output of a block
 - Allow the FP to skip certain blocks
 - Allow gradient to propagate to every layer



He et al. *Deep Residual Learning for Image Recognition*.
CVPR 2016.



Vaswani et al. *Attention Is All You Need*. NIPS 2017.

Tricks for Very Deep Models

- Batch Normalization (Layer Normalization)
 - Makes training more stable, avoid over-smoothing issues, reduce overfitting
 - <https://arxiv.org/abs/1502.03167>

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

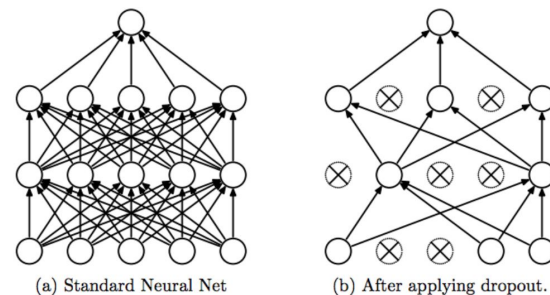
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

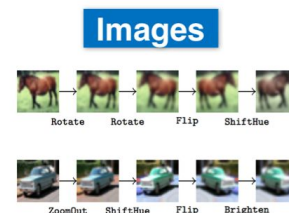
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Tricks to Reduce Overfitting

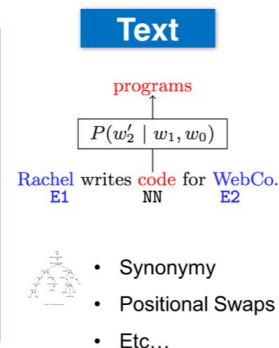
- Try overfitting the model on training set (to evaluate the expressiveness); then address the overfitting issue.
- Regularizations
 - L1/L2 regularization (weight decay): add penalty to large values in parameter
 - Dropout: disable certain output dimensions during training
- Augmentation, early stopping, etc.



<https://www.tech-quantum.com/implementing-drop-out-regularization-in-neural-networks/>



- Rotations
- Scaling / Zooms
- Brightness
- Color Shifts
- Etc...



<https://madewithml.com/courses/mlops/augmentation/>



TEXAS A&M UNIVERSITY
Engineering

Suggestions on Proposing & Building Deep Models

1. Model/method design first, hyper-tuning later

- Your model design should be based on empirical insights, reasonable motivations, theoretical frameworks, etc.
- If the method itself works, you don't even need much efforts on hyper-tuning.

2. **Start from the simplest model** when your methods consist of multiple (novel) components

- Always choose a simple base model to start with and consider it as the baseline.
- Add one component at a time to evaluate its effectiveness.

3. Reasoning the cause before making modifications when the model does not work as expected

- Observations → hypothesis → verification
 - Obtain enough evidence that gives you an idea on model behavior. E.g., visualizations, intermediate variables.
 - Make reasonable hypotheses on why models don't work.
 - Modify certain parts to verify the hypothesis.



TEXAS A&M UNIVERSITY

Engineering

Click to edit title style

Click to edit title style